Derck Prinzhorn

derckprinzhorn@gmail.com derckprinzhorn.com GitHub Scholar LinkedIn

Experience

Exponential Security Labs Jul 2026 – present
Founding Member of Technical Staff
Work on multi-agent security environments and evaluations, red-teaming frontier agents, and defenses for agentic systems.

Max Planck Institute for Intelligent Systems & ELLIS Institute Tübingen Jan – Aug 2026
Research Intern
Researched AI control with Maksym Andriushchenko, focusing on agent monitoring and oversight.

Aithos Research Lab Apr 2025 – Jan 2026
Research Engineer
Worked on evaluations of AI value systems and moral competence.

University of Amsterdam Jan 2024 – Feb 2025
Research Intern
Researched conformal prediction for time series, radiotherapy dose prediction, and physical reasoning in video generation.

Dutch Police, *AI Architect* Apr 2023 – Apr 2025

Selected Research Project

Controlling the Researcher: AI Control Evaluations for Automated AI R&D 2026
Co-developed control evaluations spanning safety fine-tuning, RL post-training, and CUDA optimization. Subtle sabotage embedded in research artifacts evaded most evaluated monitors; overt side tasks were easier to detect.

Publications

AI safety and security

Instrumental Monitor Evasion Emerges Under Ordinary Task Pressure
D. Schmotz*, **Derck Prinzhorn***, L. Beurer-Kellner, A. Paulus, A. Prabhu, M. Andriushchenko. arXiv, 2026. *Equal contribution.

LLM Agents Can Easily Tamper With Their Own Traces
J. Qin, D. Schmotz, **Derck Prinzhorn**, L. Beurer-Kellner, A. Prabhu, M. Andriushchenko. arXiv, 2026.

ResearchArena: Evaluating Sabotage and Monitoring in Automated AI R&D
L. Libon, B. Rank, J. Yeon, D. Schmotz, J. Qin, D. Donnelly, **Derck Prinzhorn**, M. Andriushchenko. COLM Workshop on Agent Behavior, 2026.

Moral Competence Before Moral Content: Why LLM Agents Lack the Prerequisites for Coherent Alignment
A. Libert, **Derck Prinzhorn**, D. R. Henselmans. Paris Journal of AI and Digital Ethics, accepted 2026.

Measuring AI Accountability Through Argumentation Analysis: Can Model Reasoning Withstand Scrutiny?
D. R. Henselmans, **Derck Prinzhorn**, A. Libert. Paris Journal of AI and Digital Ethics, accepted 2026.

Machine learning and vision

Evaluating Newtonian Mechanics in Video Generative Models with Real Physical Systems
A. Tragoudaras, C. Zhang, D. Cherniavskii, A. Vozikis, T. Nijdam, **Derck Prinzhorn**, M. Bodracska, N. Sebe, A. Zadaianchuk, E. Gavves. ICML, 2026.

Injecting Image Guidance into Text-Conditioned Diffusion Models at Inference
A. Zywot, I. Skylitsis, T. Nijdam, Z. Tzifa-Kratira, **Derck Prinzhorn**, K. Szewczyk, A. Bhowmik. arXiv, 2026.

HIVE: A Hyperbolic Interactive Visualization Explorer for Representation Learning
T. Nijdam, **Derck Prinzhorn**, J. de Heus, T. Brouwer. ICCV Workshops, 2025.

Conformal Time Series Decomposition with Component-wise Exchangeability
Derck Prinzhorn, T. Nijdam, P. A. van der Linden, A. Timans. COPA / PMLR, 2024.

Reproducibility Study of FairAC
G. de Jong, M. J. Meijer, **Derck Prinzhorn**, H. Ruiter. TMLR, 2024; NeurIPS 2024 reproducibility poster.

Education

MSc Artificial Intelligence, University of Amsterdam 2023 – 2026
Thesis research on AI control at the Max Planck Institute for Intelligent Systems and ELLIS Institute Tübingen.

BSc Artificial Intelligence, University of Amsterdam 2020 – 2023
Graduated cum laude; AmsterdamAI Thesis Award for work on uncertainty quantification in time series.